


Brain Tumor Detection using Pretrained Model using Attention Mechanism



```
import numpy as np
import pandas as pd

base_path =
"/kaggle/input/brain-tumour-classification/BrainTumor_1/Train"
categories = ["glioma", "meningioma", "notumor", "pituitary"]

image_paths = []
labels = []

for category in categories:
    category_path = os.path.join(base_path, category)
    for image_name in os.listdir(category_path):
        image_path = os.path.join(category_path, image_name)
```

```

        image_paths.append(image_path)
        labels.append(category)

df = pd.DataFrame({
    "image_path": image_paths,
    "label": labels
})

df.head()

      image_path  label
0  /kaggle/input/brain-tumour-classification/Brai...  glioma
1  /kaggle/input/brain-tumour-classification/Brai...  glioma
2  /kaggle/input/brain-tumour-classification/Brai...  glioma
3  /kaggle/input/brain-tumour-classification/Brai...  glioma
4  /kaggle/input/brain-tumour-classification/Brai...  glioma

df.tail()

      image_path  label
22843 /kaggle/input/brain-tumour-classification/Brai...  pituitary
22844 /kaggle/input/brain-tumour-classification/Brai...  pituitary
22845 /kaggle/input/brain-tumour-classification/Brai...  pituitary
22846 /kaggle/input/brain-tumour-classification/Brai...  pituitary
22847 /kaggle/input/brain-tumour-classification/Brai...  pituitary

df.shape
(22848, 2)

df.columns
Index(['image_path', 'label'], dtype='object')

df['label'].unique()
array(['glioma', 'meningioma', 'notumor', 'pituitary'], dtype=object)

df['label'].value_counts()
label
notumor      6380
pituitary    5828
meningioma   5356
glioma       5284
Name: count, dtype: int64

import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 6))
sns.countplot(data=df, x="label", palette="viridis")

```

```

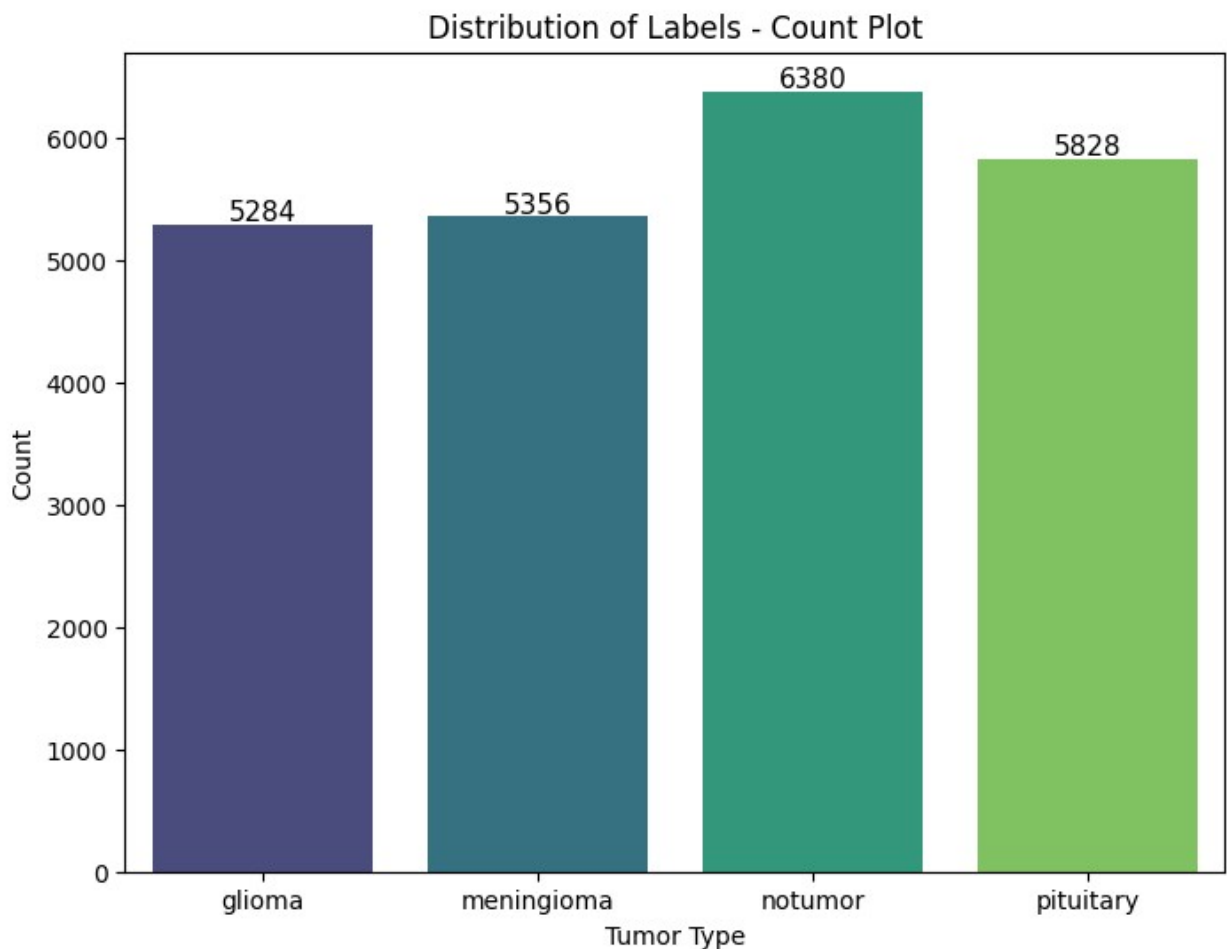
plt.title("Distribution of Labels - Count Plot")
plt.xlabel("Tumor Type")
plt.ylabel("Count")

for p in plt.gca().patches:
    plt.gca().annotate(f'{{int(p.get_height())}}',
                       (p.get_x() + p.get_width() / 2.,
                        p.get_height()),
                       ha='center', va='center', fontsize=11,
                       color='black', xytext=(0, 5),
                       textcoords='offset points')

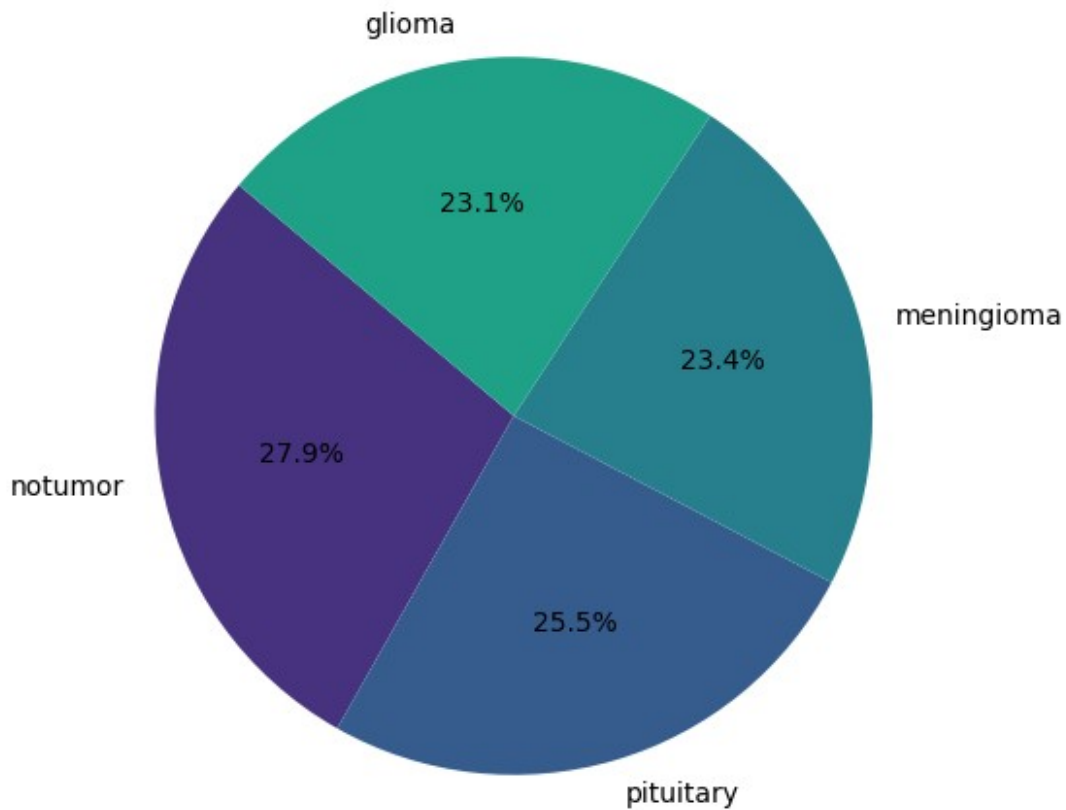
plt.show()

label_counts = df['label'].value_counts()
plt.figure(figsize=(8, 6))
plt.pie(label_counts, labels=label_counts.index, autopct='%1.1f%%',
        startangle=140, colors=sns.color_palette("viridis"))
plt.title("Distribution of Labels - Pie Chart")
plt.show()

```




Distribution of Labels - Pie Chart



```
import cv2
num_images = 5
plt.figure(figsize=(15, 12))
for i, category in enumerate(categories):
    category_images = df[df['label'] == category]
    ['image_path'].iloc[:num_images]
    for j, img_path in enumerate(category_images):
        img = cv2.imread(img_path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        plt.subplot(len(categories), num_images, i * num_images + j +
1)
        plt.imshow(img)
        plt.axis('off')
        plt.title(category)
```

```
plt.tight_layout()
plt.show()
```



```
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
df['category_encoded'] = label_encoder.fit_transform(df['label'])
df = df[['image_path', 'category_encoded']]
from imblearn.over_sampling import RandomOverSampler
ros = RandomOverSampler(random_state=42)
X_resampled, y_resampled = ros.fit_resample(df[['image_path']],
df['category_encoded'])
df_resampled = pd.DataFrame(X_resampled, columns=['image_path'])
df_resampled['category_encoded'] = y_resampled
```

```
print("\nClass distribution after oversampling:")
print(df_resampled['category_encoded'].value_counts())
```

Class distribution after oversampling:

category_encoded

0 6380

1 6380

2 6380

3 6380

Name: count, dtype: int64

df_resampled

	image_path
--	------------

category_encoded	image_path
------------------	------------

0	/kaggle/input/brain-tumour-classification/Brai...
---	---

0	
---	--

1	/kaggle/input/brain-tumour-classification/Brai...
---	---

0	
---	--

2	/kaggle/input/brain-tumour-classification/Brai...
---	---

0	
---	--

3	/kaggle/input/brain-tumour-classification/Brai...
---	---

0	
---	--

4	/kaggle/input/brain-tumour-classification/Brai...
---	---

0	
---	--

...	...
-----	-----

...	
-----	--

25515	/kaggle/input/brain-tumour-classification/Brai...
-------	---

3	
---	--

25516	/kaggle/input/brain-tumour-classification/Brai...
-------	---

3	
---	--

25517	/kaggle/input/brain-tumour-classification/Brai...
-------	---

3	
---	--

25518	/kaggle/input/brain-tumour-classification/Brai...
-------	---

3	
---	--

25519	/kaggle/input/brain-tumour-classification/Brai...
-------	---

3	
---	--

[25520 rows x 2 columns]

```
import time
import shutil
import pathlib
import itertools
from PIL import Image
```

```
import cv2
import seaborn as sns
sns.set_style('darkgrid')
```

```

import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam, Adamax
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense, Activation, Dropout, BatchNormalization
from tensorflow.keras import regularizers

import warnings
warnings.filterwarnings("ignore")

print ('check')

check

df_resampled['category_encoded'] =
df_resampled['category_encoded'].astype(str)

train_df_new, temp_df_new = train_test_split(
    df_resampled,
    train_size=0.8,
    shuffle=True,
    random_state=42,
    stratify=df_resampled['category_encoded']
)

valid_df_new, test_df_new = train_test_split(
    temp_df_new,
    test_size=0.5,
    shuffle=True,
    random_state=42,
    stratify=temp_df_new['category_encoded']
)

from tensorflow.keras.preprocessing.image import ImageDataGenerator

batch_size = 16
img_size = (224, 224)
channels = 3
img_shape = (img_size[0], img_size[1], channels)

tr_gen = ImageDataGenerator(rescale=1./255)
ts_gen = ImageDataGenerator(rescale=1./255)

train_gen_new = tr_gen.flow_from_dataframe(
    train_df_new,

```



```

from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import (GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization,
                                GaussianNoise, Input,
MultiHeadAttention, Reshape)
from tensorflow.keras.optimizers import Adam

def create_vgg16_model(input_shape):

    inputs = Input(shape=input_shape)

    base_model = VGG16(weights='imagenet', input_tensor=inputs,
include_top=False)

    for layer in base_model.layers:
        layer.trainable = False

    x = base_model.output

    height, width, channels = 7, 7, 512
    x = Reshape((height * width, channels))(x)

    attention_output = MultiHeadAttention(num_heads=8,
key_dim=channels)(x, x)
    attention_output = Reshape((height, width, channels))
(attention_output)

    x = GaussianNoise(0.25)(attention_output)

    x = GlobalAveragePooling2D()(x)

    x = Dense(512, activation='relu')(x)
    x = BatchNormalization()(x)
    x = GaussianNoise(0.25)(x)
    x = Dropout(0.25)(x)

    outputs = Dense(4, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    return model

input_shape = (224, 224, 3)

cnn_model = create_vgg16_model(input_shape)

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

```

```
history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=5,
    callbacks=[early_stopping],
    verbose=1
)
```

Epoch 1/5

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1730616392.742667 299 service.cc:145] XLA service 0x7a917c0018a0 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

I0000 00:00:1730616392.742725 299 service.cc:153] StreamExecutor device (0): Tesla T4, Compute Capability 7.5

I0000 00:00:1730616392.742730 299 service.cc:153] StreamExecutor device (1): Tesla T4, Compute Capability 7.5

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1730616404.835907 338 asm_compiler.cc:369] ptxas

warning : Registers are spilled to local memory in function 'triton_gemm_dot_1', 1284 bytes spill stores, 1272 bytes spill loads

1/1276 _____ 9:55:19 28s/step - accuracy: 0.2500 - loss: 1.9726

I0000 00:00:1730616416.859234 299 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

1276/1276 _____ 0s 78ms/step - accuracy: 0.7626 - loss: 0.6298

I0000 00:00:1730616537.083262 399 asm_compiler.cc:369] ptxas

warning : Registers are spilled to local memory in function 'triton_gemm_dot_6', 1280 bytes spill stores, 1268 bytes spill loads

1276/1276 _____ 154s 99ms/step - accuracy: 0.7627 - loss: 0.6296 - val_accuracy: 0.6759 - val_loss: 1.0673

Epoch 2/5

1276/1276 _____ 116s 90ms/step - accuracy: 0.8798 - loss: 0.3314 - val_accuracy: 0.8973 - val_loss: 0.2846

Epoch 3/5

1276/1276 _____ 116s 90ms/step - accuracy: 0.9128 - loss: 0.2491 - val_accuracy: 0.8997 - val_loss: 0.2999


Epoch 4/5


1276/1276 _____ 116s 90ms/step - accuracy: 0.9180 -

```
loss: 0.2193 - val_accuracy: 0.7022 - val_loss: 1.3966
Epoch 5/5
1276/1276 _____ 116s 90ms/step - accuracy: 0.9317 -
loss: 0.1875 - val_accuracy: 0.5247 - val_loss: 2.6919
```

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_classes = np.argmax(predictions, axis=1)
```


160/160 ————— 13s 76ms/step

```
report = classification_report(test_labels, predicted_classes,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
0	0.88	0.93	0.90	638
1	0.82	0.87	0.85	638
2	0.95	0.98	0.97	638
3	0.99	0.85	0.91	638
accuracy			0.91	2552
macro avg	0.91	0.91	0.91	2552
weighted avg	0.91	0.91	0.91	2552

```
conf_matrix = confusion_matrix(test_labels, predicted_classes)
```

```
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(test_gen_new.class_indices.keys()),
            yticklabels=list(test_gen_new.class_indices.keys()))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
from tensorflow.keras.applications import VGG19
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise, Input, MultiHeadAttention,
Reshape
from tensorflow.keras.optimizers import Adam
```

```

import tensorflow as tf

def create_vgg19_model(input_shape):
    inputs = Input(shape=input_shape)

    base_model = VGG19(weights='imagenet', input_tensor=inputs,
include_top=False)

    for layer in base_model.layers:
        layer.trainable = False

    x = base_model.output

    height, width, channels = 7, 7, 512
    x = Reshape((height * width, channels))(x)

    attention_output = MultiHeadAttention(num_heads=8,
key_dim=channels)(x, x)

    attention_output = Reshape((height, width, channels))
(attention_output)

    x = GaussianNoise(0.25)(attention_output)

    x = GlobalAveragePooling2D()(x)

    x = Dense(512, activation='relu')(x)
    x = BatchNormalization()(x)
    x = GaussianNoise(0.25)(x)
    x = Dropout(0.25)(x)

    outputs = Dense(4, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    return model

input_shape = (224, 224, 3)

cnn_model = create_vgg19_model(input_shape)

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5
80134624/80134624 ————— 0s 0us/step

```

```


history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=5,
    callbacks=[early_stopping],
    verbose=1
)

Epoch 1/5
1276/1276 _____ 149s 110ms/step - accuracy: 0.7411 -
loss: 0.6771 - val_accuracy: 0.7786 - val_loss: 0.6523
Epoch 2/5
1276/1276 _____ 138s 108ms/step - accuracy: 0.8712 -
loss: 0.3527 - val_accuracy: 0.8288 - val_loss: 0.5774
Epoch 3/5
1276/1276 _____ 138s 108ms/step - accuracy: 0.9022 -
loss: 0.2665 - val_accuracy: 0.8546 - val_loss: 0.3942
Epoch 4/5
1276/1276 _____ 138s 108ms/step - accuracy: 0.9215 -
loss: 0.2168 - val_accuracy: 0.8064 - val_loss: 0.5658
Epoch 5/5
1276/1276 _____ 138s 108ms/step - accuracy: 0.9319 -
loss: 0.1890 - val_accuracy: 0.9150 - val_loss: 0.2445

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_classes = np.argmax(predictions, axis=1)
```


160/160 ————— 15s 91ms/step

```
report = classification_report(test_labels, predicted_classes,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
0	0.98	0.81	0.89	638
1	0.83	0.92	0.87	638
2	0.97	0.97	0.97	638
3	0.92	0.98	0.95	638
accuracy			0.92	2552
macro avg	0.92	0.92	0.92	2552
weighted avg	0.92	0.92	0.92	2552

```
conf_matrix = confusion_matrix(test_labels, predicted_classes)
```

```
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(test_gen_new.class_indices.keys()),
            yticklabels=list(test_gen_new.class_indices.keys()))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
from tensorflow.keras.applications import MobileNet
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise, Input, MultiHeadAttention,
Reshape
from tensorflow.keras.optimizers import Adam
```

```

import tensorflow as tf

def create_mobilenet_model(input_shape):
    inputs = Input(shape=input_shape)

    base_model = MobileNet(weights='imagenet', input_tensor=inputs,
include_top=False)

    for layer in base_model.layers:
        layer.trainable = False

    x = base_model.output

    height, width, channels = 7, 7, 1024
    x = Reshape((height * width, channels))(x)

    attention_output = MultiHeadAttention(num_heads=8,
key_dim=channels)(x, x)

    attention_output = Reshape((height, width, channels))
(attention_output)

    x = GaussianNoise(0.25)(attention_output)

    x = GlobalAveragePooling2D()(x)

    x = Dense(512, activation='relu')(x)
    x = BatchNormalization()(x)
    x = GaussianNoise(0.25)(x)
    x = Dropout(0.25)(x)

    outputs = Dense(4, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    return model

input_shape = (224, 224, 3)

cnn_model = create_mobilenet_model(input_shape)

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet/mobilenet_1_0_224_tf_no_top.h5
 17225924/17225924 ————— 0s 0us/step

```

history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=5,
    callbacks=[early_stopping],
    verbose=1
)

```

Epoch 1/5

```

I0000 00:00:1730617968.755983      673 asm_compiler.cc:369] ptxas
warning : Registers are spilled to local memory in function
'triton_gemm_dot_1', 1300 bytes spill stores, 1284 bytes spill loads

```

```

1276/1276 _____ 0s 55ms/step - accuracy: 0.8556 - loss:
0.4049

```

```

I0000 00:00:1730618059.939100      726 asm_compiler.cc:369] ptxas
warning : Registers are spilled to local memory in function
'triton_gemm_dot_1', 1300 bytes spill stores, 1284 bytes spill loads

```

```

1276/1276 _____ 109s 68ms/step - accuracy: 0.8557 -
loss: 0.4048 - val_accuracy: 0.9412 - val_loss: 0.1564

```

Epoch 2/5

```

1276/1276 _____ 77s 60ms/step - accuracy: 0.9400 -
loss: 0.1708 - val_accuracy: 0.8374 - val_loss: 0.5594

```

Epoch 3/5

```

1276/1276 _____ 76s 59ms/step - accuracy: 0.9632 -
loss: 0.1061 - val_accuracy: 0.9208 - val_loss: 0.2246

```

Epoch 4/5

```

1276/1276 _____ 76s 59ms/step - accuracy: 0.9756 -
loss: 0.0739 - val_accuracy: 0.9671 - val_loss: 0.1134

```

Epoch 5/5

```

1276/1276 _____ 76s 60ms/step - accuracy: 0.9779 -
loss: 0.0591 - val_accuracy: 0.9659 - val_loss: 0.1115

```

```

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()


```


```

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')

```

```
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()
```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_classes = np.argmax(predictions, axis=1)
```


160/160 ————— 10s 52ms/step

```
report = classification_report(test_labels, predicted_classes,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
0	0.98	0.97	0.98	638
1	0.93	0.97	0.95	638
2	1.00	0.98	0.99	638
3	0.98	0.98	0.98	638
accuracy			0.97	2552
macro avg	0.97	0.97	0.97	2552
weighted avg	0.97	0.97	0.97	2552

```
conf_matrix = confusion_matrix(test_labels, predicted_classes)
```

```
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(test_gen_new.class_indices.keys()),
            yticklabels=list(test_gen_new.class_indices.keys()))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
from tensorflow.keras.applications import Xception
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise, Input, MultiHeadAttention,
Reshape
from tensorflow.keras.optimizers import Adam
```

```

import tensorflow as tf

def create_xception_model(input_shape):
    inputs = Input(shape=input_shape)

    base_model = Xception(weights='imagenet', input_tensor=inputs,
include_top=False)

    for layer in base_model.layers:
        layer.trainable = False

    x = base_model.output

    height, width, channels = 7, 7, 2048
    x = Reshape((height * width, channels))(x)

    attention_output = MultiHeadAttention(num_heads=8,
key_dim=channels)(x, x)

    attention_output = Reshape((height, width, channels))
(attention_output)

    x = GaussianNoise(0.25)(attention_output)
    x = GlobalAveragePooling2D()(x)
    x = Dense(512, activation='relu')(x)
    x = BatchNormalization()(x)
    x = GaussianNoise(0.25)(x)
    x = Dropout(0.25)(x)
    outputs = Dense(4, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    return model

input_shape = (224, 224, 3)
cnn_model = create_xception_model(input_shape)

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/xception/
xception_weights_tf_dim_ordering_tf_kernels_notop.h5
83683744/83683744 _____ 0s 0us/step

history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,

```



```

    epochs=5,
    callbacks=[early_stopping],
    verbose=1
)

Epoch 1/5
I0000 00:00:1730618526.007229      855 asm_compiler.cc:369] ptxas
warning : Registers are spilled to local memory in function
'triton_gemm_dot_1', 1300 bytes spill stores, 1284 bytes spill loads

1276/1276 _____ 0s 231ms/step - accuracy: 0.7874 -
loss: 0.5876


I0000 00:00:1730618864.402637      915 asm_compiler.cc:369] ptxas
warning : Registers are spilled to local memory in function
'triton_gemm_dot_1', 1300 bytes spill stores, 1284 bytes spill loads


1276/1276 _____ 362s 260ms/step - accuracy: 0.7874 -
loss: 0.5875 - val_accuracy: 0.9087 - val_loss: 0.2856
Epoch 2/5
1276/1276 _____ 317s 248ms/step - accuracy: 0.8945 -
loss: 0.2957 - val_accuracy: 0.9032 - val_loss: 0.2573
Epoch 3/5
1276/1276 _____ 315s 247ms/step - accuracy: 0.9206 -
loss: 0.2221 - val_accuracy: 0.8997 - val_loss: 0.2763
Epoch 4/5
1276/1276 _____ 315s 247ms/step - accuracy: 0.9414 -
loss: 0.1703 - val_accuracy: 0.9487 - val_loss: 0.1454
Epoch 5/5
1276/1276 _____ 315s 247ms/step - accuracy: 0.9532 -
loss: 0.1366 - val_accuracy: 0.9549 - val_loss: 0.1152

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_classes = np.argmax(predictions, axis=1)
```


160/160 ————— 25s 139ms/step

```
report = classification_report(test_labels, predicted_classes,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
0	0.91	0.94	0.92	638
1	0.87	0.80	0.83	638
2	0.90	0.98	0.94	638
3	0.95	0.90	0.92	638
accuracy			0.91	2552
macro avg	0.91	0.91	0.90	2552
weighted avg	0.91	0.91	0.90	2552

```
conf_matrix = confusion_matrix(test_labels, predicted_classes)
```

```
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(test_gen_new.class_indices.keys()),
            yticklabels=list(test_gen_new.class_indices.keys()))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
from tensorflow.keras.applications import InceptionV3
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise, Input, MultiHeadAttention,
Reshape
from tensorflow.keras.optimizers import Adam
```

```

import tensorflow as tf

def create_inception_model(input_shape):
    inputs = Input(shape=input_shape)

    base_model = InceptionV3(weights='imagenet', input_tensor=inputs,
include_top=False)

    for layer in base_model.layers:
        layer.trainable = False

    x = base_model.output

    height, width, channels = 5, 5, 2048
    x = Reshape((height * width, channels))(x)

    attention_output = MultiHeadAttention(num_heads=8,
key_dim=channels)(x, x)

    attention_output = Reshape((height, width, channels))
(attention_output)

    x = GaussianNoise(0.25)(attention_output)
    x = GlobalAveragePooling2D()(x)
    x = Dense(512, activation='relu')(x)
    x = BatchNormalization()(x)
    x = GaussianNoise(0.25)(x)
    x = Dropout(0.25)(x)
    outputs = Dense(4, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    return model

input_shape = (224, 224, 3)
cnn_model = create_inception_model(input_shape)

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/inception_v3/
inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5
87910968/87910968 ————— 0s 0us/step

history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=5,

```

```

        callbacks=[early_stopping],
        verbose=1
    )

Epoch 1/5

I0000 00:00:1730620812.574405    1045 asm_compiler.cc:369] ptxas
warning : Registers are spilled to local memory in function
'triton_gemm_dot_1', 1300 bytes spill stores, 1284 bytes spill loads

1276/1276 _____ 0s 127ms/step - accuracy: 0.7763 -
loss: 0.6269

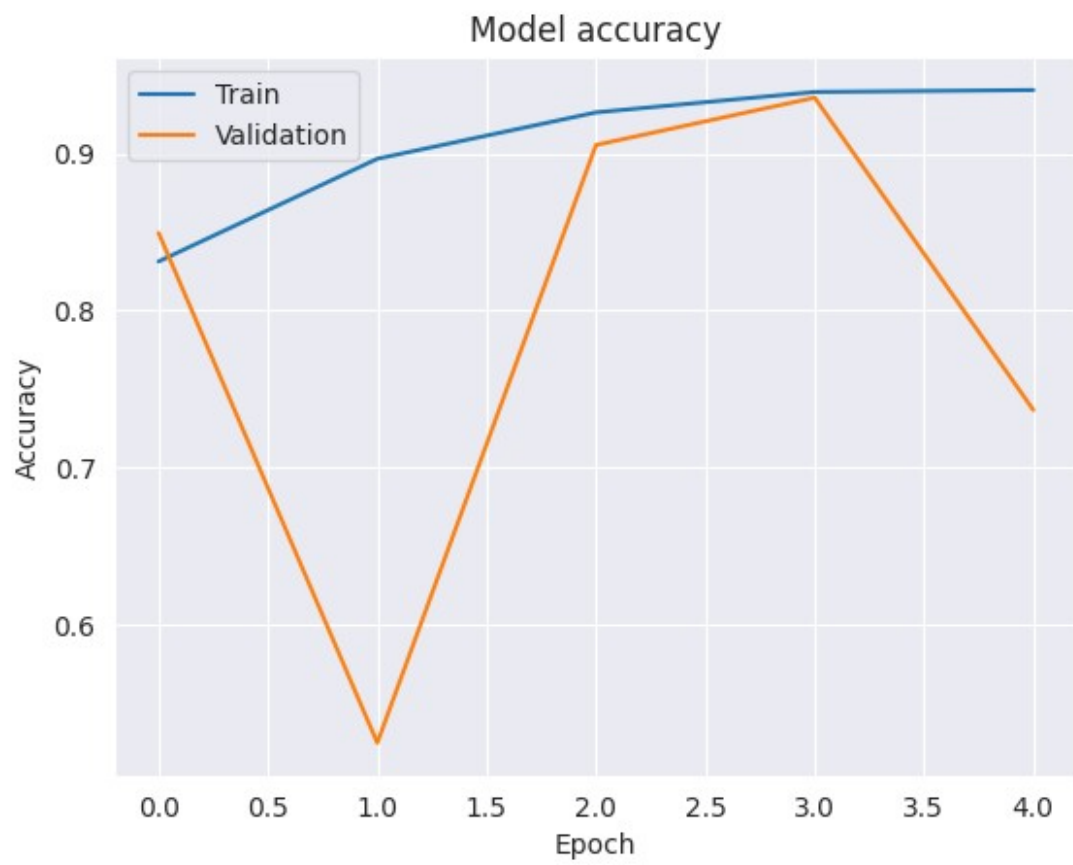
I0000 00:00:1730621007.694029    1096 asm_compiler.cc:369] ptxas
warning : Registers are spilled to local memory in function
'triton_gemm_dot_1', 1300 bytes spill stores, 1284 bytes spill loads


1276/1276 _____ 221s 149ms/step - accuracy: 0.7764 -
loss: 0.6267 - val_accuracy: 0.8495 - val_loss: 0.5046
Epoch 2/5
1276/1276 _____ 172s 135ms/step - accuracy: 0.8994 -
loss: 0.2807 - val_accuracy: 0.5247 - val_loss: 1.6268
Epoch 3/5
1276/1276 _____ 171s 134ms/step - accuracy: 0.9281 -
loss: 0.2047 - val_accuracy: 0.9056 - val_loss: 0.2910
Epoch 4/5
1276/1276 _____ 171s 134ms/step - accuracy: 0.9375 -
loss: 0.1739 - val_accuracy: 0.9357 - val_loss: 0.1815
Epoch 5/5
1276/1276 _____ 172s 134ms/step - accuracy: 0.9478 -
loss: 0.1452 - val_accuracy: 0.7367 - val_loss: 0.9598

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_classes = np.argmax(predictions, axis=1)
```

160/160 ————— 20s 94ms/step


```
report = classification_report(test_labels, predicted_classes,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
0	0.91	0.90	0.90	638
1	0.93	0.53	0.68	638
2	0.94	0.96	0.95	638
3	0.71	1.00	0.83	638
accuracy			0.85	2552
macro avg	0.87	0.85	0.84	2552
weighted avg	0.87	0.85	0.84	2552

```
conf_matrix = confusion_matrix(test_labels, predicted_classes)
```



```
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(test_gen_new.class_indices.keys()),
            yticklabels=list(test_gen_new.class_indices.keys()))
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
models = ['VGG16', 'VGG19', 'MobileNet', 'Xception', 'InceptionV3']
precision_scores = [0.91, 0.92, 0.97, 0.91, 0.87]
recall_scores = [0.91, 0.92, 0.97, 0.91, 0.85]
f1_scores = [0.91, 0.92, 0.97, 0.90, 0.84]
accuracy_scores = [0.91, 0.92, 0.97, 0.91, 0.85]

bar_width = 0.2
```

```

index = np.arange(len(models))

fig, ax = plt.subplots(figsize=(12, 8))

bars1 = ax.bar(index, precision_scores, bar_width, label='Precision')
bars2 = ax.bar(index + bar_width, recall_scores, bar_width,
label='Recall')
bars3 = ax.bar(index + 2 * bar_width, f1_scores, bar_width, label='F1
Score')
bars4 = ax.bar(index + 3 * bar_width, accuracy_scores, bar_width,
label='Accuracy')


def add_labels(bars):
    for bar in bars:
        height = bar.get_height()
        ax.annotate(f'{height:.2f}', xy=(bar.get_x() + bar.get_width()
/ 2, height),
                    xytext=(0, 3), textcoords="offset points",
ha='center', va='bottom')

add_labels(bars1)
add_labels(bars2)
add_labels(bars3)
add_labels(bars4)

ax.set_xlabel('Models')
ax.set_ylabel('Scores')
ax.set_title('Comparative Metrics for Different Models')
ax.set_xticks(index + bar_width * 1.5)
ax.set_xticklabels(models)
ax.legend()

plt.tight_layout()
plt.show()

```



Copyright @ThinkAI - A Machine Learning Community